

Algorithms at Scale

(Week 13)

Dimensionality Reduction

Giovanni
Pagliarini

Corentin
Dumery


Summary

- Motivation
- Approaches
- Johnson-Lindenstrauss Transforms
- Fast JL Transform (+ Proof)
- Experiments' Design
- Results

Dimensionality Reduction - Why?

				d
n	0.3	0.9	...	0.343
	0.9	0.2	...	0.43

	0.6	0.0	...	1.0



- Reduces cpu workload and storage needed;
- Alleviates curse of dimensionality;
- Allows noise/redundancies removal.

} potential
performance
improvement

Common approaches

- Drop useless variables
 - Variables with **low variance** may not tell anything useful;
 - When variables are **highly correlated**, one of them is unnecessary;
- Derive a new set of variables
 - Factor Analysis;
 - Principal Component Analysis;
 - Linear Discriminant Analysis;
 - ...
 - **Randomized Projections**

d			
0.3	0.9	...	0.343
0.9	0.2	...	0.43
...
0.6	0.0	...	1.0

n

What do we want?

A function that **reduces dimensionality**:

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

What do we want?

A function that **reduces dimensionality**:

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

Without **losing information**:

$$\forall u, v \in X \subset \mathbb{R}^d, (1 - \epsilon)\|u - v\|^2 \leq \|f(u) - f(v)\|^2 \leq (1 + \epsilon)\|u - v\|^2$$

What do we want?

A function that **reduces dimensionality**:

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

Without **losing information**:

$$\forall u, v \in X \subset \mathbb{R}^d, (1 - \epsilon)\|u - v\|^2 \leq \|f(u) - f(v)\|^2 \leq (1 + \epsilon)\|u - v\|^2$$

Johnson-Lindenstrauss Lemma:

Such a function exists.

[1] W.B. Johnson, J. Lindenstrauss, Extensions of Lipschitz mappings into a Hilbert space, Conference in modern analysis and probability, New Haven, CT, 1982, Amer. Math. Soc., Providence, RI, 1984, pp. 189-206

Can we build a JL Transform?

- We call **Johnson-Lindenstrauss transform** any function that satisfies these properties.
- Recall that in the lemma f has to satisfy:

$$\forall u, v \in X \subset \mathbb{R}^d, (1 - \epsilon)\|u - v\|^2 \leq \|f(u) - f(v)\|^2 \leq (1 + \epsilon)\|u - v\|^2$$

Can we build a JL Transform?

- We call **Johnson-Lindenstrauss transform** any function that satisfies these properties.
- Recall that in the lemma f has to satisfy:

$$\forall u, v \in X \subset \mathbb{R}^d, (1 - \epsilon)\|u - v\|^2 \leq \|f(u) - f(v)\|^2 \leq (1 + \epsilon)\|u - v\|^2$$

f depends on X !

□ Our data is too big, we can't go through the whole input data to define f ...

Let's tweak our definition to make it easier.

What do we want?

A function that **reduces dimensionality**:

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

Without **losing information** **with probability at least $\frac{2}{3}$** :

$$\forall u, v \in X \subset \mathbb{R}^d, (1 - \epsilon)\|u - v\|^2 \leq \|f(u) - f(v)\|^2 \leq (1 + \epsilon)\|u - v\|^2$$

Johnson-Lindenstrauss Lemma:

Such a function exists.

[1] W.B. Johnson, J. Lindenstrauss, Extensions of Lipschitz mappings into a Hilbert space, Conference in modern analysis and probability, New Haven, CT, 1982, Amer. Math. Soc., Providence, RI, 1984, pp. 189-206

What's the simplest way of making new dimensions you can think of?

Binary coins


- Just take random linear combinations of the original dimensions.
- Very simple to implement:

buildTransform() :

```
P is a  $d \times n$  matrix  
pick  $P[i][j]$  in  $\{-1, 0, 1\}$   
return P
```

transform(vector) :

```
return P*vector
```



[2] D. Achlioptas. Database-friendly random projections: Johnson-Lindenstrauss with binary coins. J. Comput. Syst. Sci., 66(4):671-687, 2003.

Matrix multiplications can be slow.
How can we make this faster ?

The Sparsity principle

- If the matrix is sparse, the multiplication is faster to compute.

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 22 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 17 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 41 & 0 & 0 \end{bmatrix}$$

The Sparsity principle

- If the matrix is sparse, the multiplication is faster to compute.

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 22 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 17 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 41 & 0 & 0 \end{bmatrix}$$


□ Can simply be stored as :

- 22 in (1,1)
- 17 in (2,4)
- 41 in (4,3)

□ If the fraction of non-zero elements is $1/k$, then the multiplication is (approximately) k times faster !


The Sparsity principle

- If the matrix is sparse, the multiplication is faster to compute.



- If the coin-flip is very likely to fail, the transform matrix will be sparse and the operation will be fast.

A problem



What if we miss the 1 in all of our new dimensions ?

Fast JL Transform (FJLT)

- We want to pre-process the data to avoid the previous problem (with high probability).


New transform:

$$f(x) = PHDx$$

P is
sparse

$u = HDx$ is the
pre-processed vector

How do we design H
and D?



[3] Ailon, N., Chazelle, B.: The fast Johnson-Lindenstrauss transform and approximate nearest neighbors. SIAM J. Comput. 39(1), 302-322 (2009)

FJLT: $f(x) = PHDx$

- P is the same kind of coin-flip matrix but with a very biased coin
-
-

FJLT: $f(x) = PHDx$

- P : coin-flip matrix

- $H = d^{-1/2} \begin{bmatrix} +1 & +1 & +1 & +1 & +1 & +1 & +1 & +1 \\ +1 & -1 & +1 & -1 & +1 & -1 & +1 & -1 \\ +1 & +1 & -1 & -1 & +1 & +1 & -1 & -1 \\ +1 & -1 & -1 & +1 & -1 & -1 & +1 & +1 \\ +1 & +1 & +1 & +1 & -1 & -1 & -1 & -1 \\ +1 & -1 & +1 & -1 & -1 & +1 & -1 & +1 \\ +1 & +1 & -1 & -1 & -1 & -1 & +1 & +1 \\ +1 & -1 & -1 & +1 & -1 & +1 & +1 & -1 \end{bmatrix}$

-

(H is a Walsh-Hadamard matrix.

More formally, H is a square matrix of size d and: $H_{i,j} = d^{-1/2}(-1)^{\langle i-1, j-1 \rangle}$)

FJLT: $f(x) = PHDx$

- P : coin-flip matrix

- $H = d^{-1/2} \begin{bmatrix} +1 & +1 & +1 & +1 & +1 & +1 & +1 & +1 \\ +1 & -1 & +1 & -1 & +1 & -1 & +1 & -1 \\ +1 & +1 & -1 & -1 & +1 & +1 & -1 & -1 \\ +1 & -1 & -1 & +1 & -1 & -1 & +1 & +1 \\ +1 & +1 & +1 & +1 & -1 & -1 & -1 & -1 \\ +1 & -1 & +1 & -1 & -1 & +1 & -1 & +1 \\ +1 & +1 & -1 & -1 & -1 & -1 & +1 & +1 \\ +1 & -1 & -1 & +1 & -1 & +1 & +1 & -1 \end{bmatrix}$

(H is a Walsh-Hadamard matrix.

More formally, H is a square matrix of size d and: $H_{i,j} = d^{-1/2}(-1)^{\langle i-1, j-1 \rangle}$)

□ H spreads out the vector:

- $\begin{bmatrix} +1 & +1 & +1 \\ +1 & -1 & +1 \\ +1 & +1 & -1 \end{bmatrix} \begin{bmatrix} \text{red} \\ \text{blue} \\ \text{green} \end{bmatrix} = \begin{bmatrix} \text{red} + \text{blue} + \text{green} \\ \text{red} - \text{blue} + \text{green} \\ \text{red} + \text{blue} - \text{green} \end{bmatrix}$

FJLT: $f(x) = PHDx$

- P : coin-flip matrix


- $H = d^{-1/2} \begin{bmatrix} +1 & +1 & +1 & +1 & +1 & +1 & +1 & +1 \\ +1 & -1 & +1 & -1 & +1 & -1 & +1 & -1 \\ +1 & +1 & -1 & -1 & +1 & +1 & -1 & -1 \\ +1 & -1 & -1 & +1 & -1 & -1 & +1 & +1 \\ +1 & +1 & +1 & +1 & -1 & -1 & -1 & -1 \\ +1 & -1 & +1 & -1 & -1 & +1 & -1 & +1 \\ +1 & +1 & -1 & -1 & -1 & -1 & +1 & +1 \\ +1 & -1 & -1 & +1 & -1 & +1 & +1 & -1 \end{bmatrix}$

-

(H is a Walsh-Hadamard matrix.

More formally, H is a square matrix of size d and: $H_{i,j} = d^{-1/2}(-1)^{\langle i-1, j-1 \rangle}$)

□ This will solve our problems



FJLT: $f(x) = PHDx$

- P : coin-flip matrix

- $H = d^{-1/2} \begin{bmatrix} +1 & +1 & +1 & +1 & +1 & +1 & +1 & +1 \\ +1 & -1 & +1 & -1 & +1 & -1 & +1 & -1 \\ +1 & +1 & -1 & -1 & +1 & +1 & -1 & -1 \\ +1 & -1 & -1 & +1 & -1 & -1 & +1 & +1 \\ +1 & +1 & +1 & +1 & -1 & -1 & -1 & -1 \\ +1 & -1 & +1 & -1 & -1 & +1 & -1 & +1 \\ +1 & +1 & -1 & -1 & -1 & -1 & +1 & +1 \\ +1 & -1 & -1 & +1 & -1 & +1 & +1 & -1 \end{bmatrix}$

- $D = \begin{bmatrix} \pm 1 & & \\ & \ddots & \\ & & \pm 1 \end{bmatrix}$

□ Randomly switching signs of variables:
we hope the values will cancel out in the
aggregation

FJLT: $f(x) = PHDx$

- P : coin-flip matrix

← $\frac{\varepsilon^{-1} \log n}{d}$ -sparse

- $H = d^{-1/2}$

$$\begin{bmatrix} +1 & +1 & +1 & +1 & +1 & +1 & +1 & +1 \\ +1 & -1 & +1 & -1 & +1 & -1 & +1 & -1 \\ +1 & +1 & -1 & -1 & +1 & +1 & -1 & -1 \\ +1 & -1 & -1 & +1 & -1 & -1 & +1 & +1 \\ +1 & +1 & +1 & +1 & -1 & -1 & -1 & -1 \\ +1 & -1 & +1 & -1 & -1 & +1 & -1 & +1 \\ +1 & +1 & -1 & -1 & -1 & -1 & +1 & +1 \\ +1 & -1 & -1 & +1 & -1 & +1 & +1 & -1 \end{bmatrix}$$

Walsh-Hadamard is well-known : using the FFT the multiplication takes $O(d \log(d))$ time

- $D = \begin{bmatrix} \pm 1 & & \\ & \ddots & \\ & & \pm 1 \end{bmatrix}$

← $O(d)$ instead of $O(d^2)$

$$\square O(d \log(d) + d \varepsilon^{-1} \log n)$$

FJLT properties and high-level proof

Unbiased property: $(1 - \epsilon)\alpha\|x\|_1 \leq \mathbb{E}[\|y\|_1] \leq (1 + \epsilon)\alpha\|x\|_1$

where α is a constant factor. ($\alpha = k\sqrt{2\pi^{-1}}$)

FJLT properties and high-level proof

Unbiased property: $(1 - \epsilon)\alpha\|x\|_1 \leq \mathbb{E}[\|y\|_1] \leq (1 + \epsilon)\alpha\|x\|_1$

where α is a constant factor. ($\alpha = k\sqrt{2\pi^{-1}}$)

Sharp concentration around the mean property:

$$\Pr[(1 - \epsilon)\mathbb{E}[\|y\|_1] \leq \|y\|_1 \leq (1 + \epsilon)\mathbb{E}[\|y\|_1]] \geq 1 - 1/20$$

FJLT properties and high-level proof

Unbiased property: $(1 - \epsilon)\alpha\|x\|_1 \leq \mathbb{E}[\|y\|_1] \leq (1 + \epsilon)\alpha\|x\|_1$

where α is a constant factor. ($\alpha = k\sqrt{2\pi^{-1}}$)

Sharp concentration around the mean property:

$$\Pr[(1 - \epsilon)\mathbb{E}[\|y\|_1] \leq \|y\|_1 \leq (1 + \epsilon)\mathbb{E}[\|y\|_1]] \geq 1 - 1/20$$

Not-too-big lemma:

$$\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$$

FJLT properties and high-level proof

Unbiased property: $(1 - \epsilon)\alpha\|x\|_1 \leq \mathbb{E}[\|y\|_1] \leq (1 + \epsilon)\alpha\|x\|_1$

where α is a constant factor. ($\alpha = k\sqrt{2\pi^{-1}}$)

Sharp concentration around the mean property:

$$\Pr[(1 - \epsilon)\mathbb{E}[\|y\|_1] \leq \|y\|_1 \leq (1 + \epsilon)\mathbb{E}[\|y\|_1]] \geq 1 - 1/20$$

Not-too-big lemma: pre-processed data

$$\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$$



FJLT: $f(x) = PHDx$


- P : coin-flip matrix

- $H = d^{-1/2} \begin{bmatrix} +1 & +1 & +1 & +1 & +1 & +1 & +1 & +1 \\ +1 & -1 & +1 & -1 & +1 & -1 & +1 & -1 \\ +1 & +1 & -1 & -1 & +1 & +1 & -1 & -1 \\ +1 & -1 & -1 & +1 & -1 & -1 & +1 & +1 \\ +1 & +1 & +1 & +1 & -1 & -1 & -1 & -1 \\ +1 & -1 & +1 & -1 & -1 & +1 & -1 & +1 \\ +1 & +1 & -1 & -1 & -1 & -1 & +1 & +1 \\ +1 & -1 & -1 & +1 & -1 & +1 & +1 & -1 \end{bmatrix}$

□ Problem : H may actually produce much larger values

- $D = \begin{bmatrix} \pm 1 & & & \\ & \ddots & & \\ & & \ddots & \\ & & & \pm 1 \end{bmatrix}$

Another problem



What happens if we aggregate a big value with small ones ?

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

By construction: $u_1 = \sum_{i=1}^d \pm d^{-1/2} x_i$

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$

$$H_{i,j} = d^{-1/2} (-1)^{\langle 1-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

By construction: $u_1 = \sum_{i=1}^d \pm d^{-1/2} x_i$

$E[e^{sdu_1}]$ ← Moment generating function

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$

$$H_{i,j} = d^{-1/2} (-1)^{\langle 1-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

By construction: $u_1 = \sum_{i=1}^d \pm d^{-1/2} x_i$

$$\begin{aligned} \mathbb{E}[e^{sdu_1}] &= \prod_{i=1}^d \mathbb{E}[e^{sd(\pm d^{-1/2})x_i}] \\ &= \prod_{i=1}^d \left(\frac{1}{2} e^{s\sqrt{d}x_i} + \frac{1}{2} e^{-s\sqrt{d}x_i} \right) \end{aligned}$$

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$

$$H_{i,j} = d^{-1/2} (-1)^{\langle 1-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

By construction: $u_1 = \sum_{i=1}^d \pm d^{-1/2} x_i$

$$\mathbb{E}[e^{sdu_1}] = \prod_{i=1}^d \mathbb{E}[e^{sd(\pm d^{-1/2})x_i}]$$

$$= \prod_{i=1}^d \cosh(s\sqrt{d}x_i)$$

$$\leq \prod_{i=1}^d e^{s^2 dx_i^2/2} \quad (\text{using } \cosh(x) \leq e^{x^2/2})$$

$$\leq e^{s^2 d \|x\|_2^2/2}$$

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$

$$H_{i,j} = d^{-1/2} (-1)^{\langle 1-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

$$\mathbb{E}[e^{sdu_1}] \leq e^{s^2 d \|x\|_2^2/2}$$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

$$\Pr[|u_1| \geq s] = \Pr[u_1 \geq s] + \Pr[u_1 \leq -s]$$

By symmetry: $\Pr[|u_1| \geq s] = 2 \Pr[e^{sdu_1} \geq e^{s^2d}]$

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$

$$H_{i,j} = d^{-1/2} (-1)^{\langle i-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

$$\mathbb{E}[e^{sdu_1}] \leq e^{s^2d\|x\|_2^2/2}$$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

$$\Pr[|u_1| \geq s] = \Pr[u_1 \geq s] + \Pr[u_1 \leq -s]$$

By symmetry: $\Pr[|u_1| \geq s] = 2 \Pr[e^{sdu_1} \geq e^{s^2d}]$

$$\begin{aligned} \Pr[|u_1| \geq s] &= 2 \Pr[e^{sdu_1} \geq e^{s^2d}] \\ &\leq 2\mathbb{E}[e^{sdu_1}] / e^{s^2d} \end{aligned}$$

(Markov's inequality)

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$

$$H_{i,j} = d^{-1/2} (-1)^{\langle i-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

$$\mathbb{E}[e^{sdu_1}] \leq e^{s^2d\|x\|_2^2/2}$$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

$$\Pr[|u_1| \geq s] = \Pr[u_1 \geq s] + \Pr[u_1 \leq -s]$$

By symmetry: $\Pr[|u_1| \geq s] = 2 \Pr[e^{sdu_1} \geq e^{s^2d}]$

$$\Pr[|u_1| \geq s] = 2 \Pr[e^{sdu_1} \geq e^{s^2d}]$$

$$\leq 2\mathbb{E}[e^{sdu_1}] / e^{s^2d}$$

(Markov's inequality)

$$\leq 2e^{s^2d(\|x\|_2^2/2-1)}$$

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$

$$H_{i,j} = d^{-1/2}(-1)^{\langle i-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

$$\mathbb{E}[e^{sdu_1}] \leq e^{s^2d\|x\|_2^2/2}$$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

$$\Pr[|u_1| \geq s] = \Pr[u_1 \geq s] + \Pr[u_1 \leq -s]$$

By symmetry: $\Pr[|u_1| \geq s] = 2 \Pr[e^{sdu_1} \geq e^{s^2d}]$

$$\Pr[|u_1| \geq s] = 2 \Pr[e^{sdu_1} \geq e^{s^2d}]$$

$$\leq 2\mathbb{E}[e^{sdu_1}] / e^{s^2d}$$

(Markov's inequality)

$$\leq 2e^{s^2d(\|x\|_2^2/2 - 1)}$$

$$\leq 2e^{-s^2d/2} \quad (\text{assuming } \|x\|_2 = 1)$$

$$\Pr[|u_1| \geq d^{-1/2} \sqrt{\log(40n)}] \leq 2e^{-\log(40n)/2}$$

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$

$$H_{i,j} = d^{-1/2} (-1)^{\langle i-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

$$\mathbb{E}[e^{sdu_1}] \leq e^{s^2d\|x\|_2^2/2}$$

Claim: $\max_{\substack{x \in X \\ \|x\|_2 = 1}} \|HDx\|_\infty = O(d^{-1/2} \sqrt{\log(n)})$

$$\Pr[|u_1| \geq d^{-1/2} \sqrt{\log(40n)}] \leq 2e^{-\log(40n)/2} \\ \leq 1/(20nd)$$

$$\Pr[\|u\|_\infty \geq d^{-1/2} \sqrt{\log(40n)}] \leq 1/(20n)$$

$$\Pr[\exists x \in X, \|HDx\|_\infty \geq d^{-1/2} \sqrt{\log(40n)}] \leq 1/20$$

Must-remember box

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^k, k \ll d$$

$$u := HDx$$


$$H_{i,j} = d^{-1/2} (-1)^{\langle 1-1, j-1 \rangle}$$

$$D_{i,i} = \pm 1$$

$$\mathbb{E}[e^{sdu_1}] \leq e^{s^2 d \|x\|_2^2 / 2}$$



Experiments' Design

We assessed the goodness of some dim. reduction methods



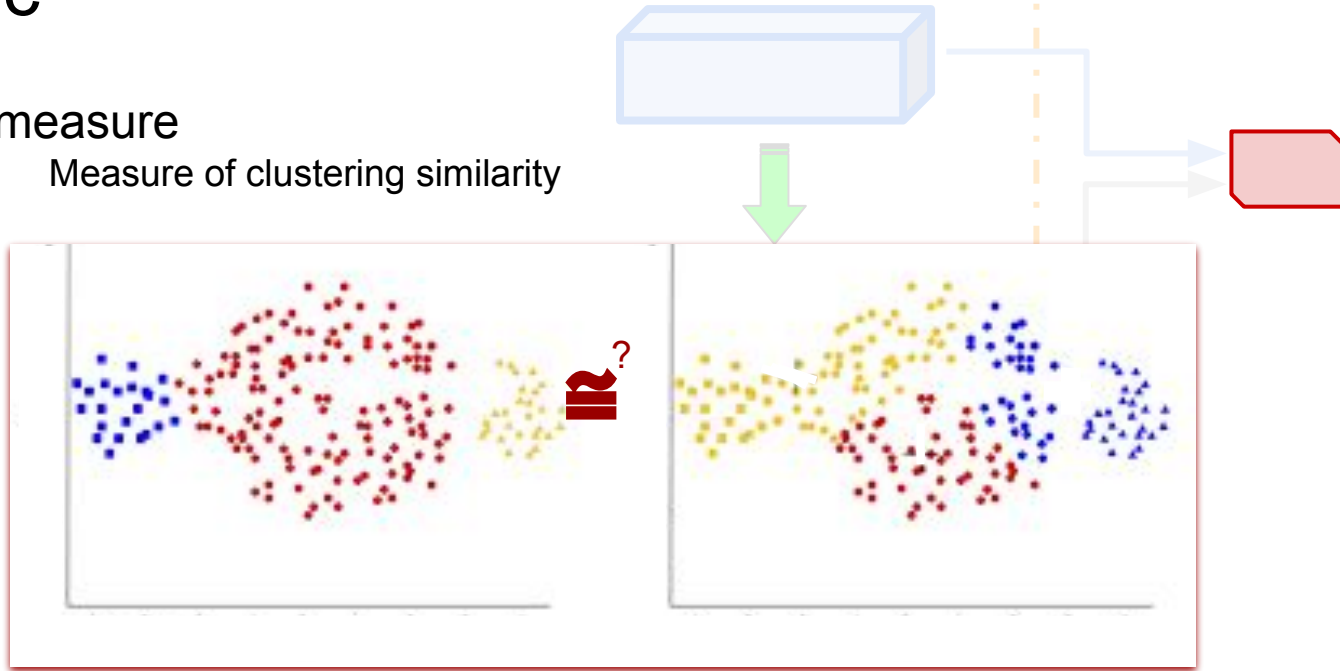
Test problems

- K-means Clustering
 - Elliptical clusters
- KNN Clustering
 - Complex shaped clusters



Metric


- v-measure
 - Measure of clustering similarity



[4] Andrew Rosenberg, Julia Hirschberg, V-Measure: A Conditional Entropy-Based External Cluster Evaluation Measure, Proceedings of the 2007 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning (EMNLP-CoNLL), 2007, pp. 410–420


Datasets

- MNIST
 - $d = 28 \times 28 = 784$
 - $N = 10'000$
 - # clusters = 10
- GISETTE
 - $d = 5'000$
 - $N = 7'000$
 - # clusters = 2
- Synthetic high-dimensional data
 - $d = \dots, 256, 512, 1024$
 - $N = 1024$
 - # clusters = 16




Methods

- Low variance filter
 - Select the k most variant variables
- Sampling variables
 - Random sample k variables
- Coin Flip method
 - Linear project on a random k -hyperplane
- Fast JL Transform
 - Random projection with improved bounds
- Principal Component Analysis
 - Presumably explained about 20 minutes ago




Result

K-Means, Synthetic data ($d = 1024$)




Result

K-Means, MNIST ($d = 784$)




Result

K-Nearest Neighbors MNIST ($d = 784$)




Result

K-Means, GISETTE ($d = 5000$)



Result

K-Means, GISETTE ($d = 5000$)



Bibliography

- [1] W.B. Johnson, J. Lindenstrauss, Extensions of Lipschitz mappings into a Hilbert space, Conference in modern analysis and probability, New Haven, CI, 1982, Amer. Math. Soc., Providence, RI, 1984, pp. 189-206
- [2] D. Achlioptas. Database-friendly random projections: Johnson-Lindenstrauss with binary coins. J. Comput. Syst. Sci., 66(4):671-687, 2003
- [3] Ailon, N., Chazelle, B.: The fast Johnson-Lindenstrauss transform and approximate nearest neighbors. SIAM J. Comput. 39(1), 302-322 (2009)
- [4] Andrew Rosenberg, Julia Hirschberg, V-Measure: A Conditional Entropy-Based External Cluster Evaluation Measure, Proceedings of the 2007 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning (EMNLP-CoNLL), 2007, pp. 410–420